

The Value of a Calibrated Error Bar Is Real Where the Parameter Is Identifiable: An End-Stage Synthesis of Trust, Value of Information, and Action in Quantitative IVIM Imaging

Avery Karlin

Department of Applied Physics and Applied Mathematics, Columbia University, New York, NY, USA

Department of Computer Science, University of Colorado Boulder, Boulder, CO, USA

The Annealing Signet Institute

ORCID: [0000-0003-3848-6782](https://orcid.org/0000-0003-3848-6782)

Abstract

A deployed quantitative-MRI parameter map carries an error bar; a clinician who acts on it faces three questions in order — can the bar be *trusted*, is a trusted bar *worth* anything to the decision, and where does acting on it *break*? This perspective synthesizes a multi-project intravoxel-incoherent-motion (IVIM) uncertainty program into one arc, trust $\rightarrow$ value of information $\rightarrow$ action, and argues that the program answers the three questions as a single chain. It makes no new measurement. The synthesis terminates in a sharp, defensible negative anchored on the pseudo-diffusion coefficient D^* : D^* is *cross-modally orphaned*. It is unidentifiable from its own signal — a Cramér–Rao identifiability bound that grows $\sim 6\times$ across the D^* range and reaches the regime-bin scale in the high- D^* tercile (segmented IVIM, SNR 10–100); it is un-scalable to its own repeatability — on real same-day test–retest data its interval width does *not* track scan–rescan scatter (Spearman $r = -0.17$, $p = 0.13$, $n = 76$; 95% bootstrap CI $[-0.39, +0.05]$, spanning zero), where the well-identified diffusion coefficient D does ($r = +0.60$); and it is only weakly and *cohort-inconsistently* tied to an independent perfusion readout (DCE K^{trans} : $r = 0.389$, $p < 0.001$ in one cohort but null in another). The end-stage caution is therefore not “uncertainty quantification fails” but the sharper claim its own machinery earns: the value of a calibrated error bar is real and quantifiable exactly where the parameter is identifiable, and D^* marks the identifiability wall where trust, value, and action all terminate together.

Provisional status. This is an end-stage synthesis assembled before its in-program anchors have published. Claims marked ^P depend on an unpublished anchor; the load-bearing spine (the D^* thread) rests on in-repository, reproduced anchors and survives regardless. Submission is held until the trust and value anchors publish (Sec. 7).

Keywords: intravoxel incoherent motion (IVIM); uncertainty quantification; value of information; identifiability; pseudo-diffusion; conformal prediction; decision-making.

1 Introduction: three questions a clinician actually asks

Quantitative MRI increasingly ships not just a parameter map but an *error bar* on every voxel. The premise of uncertainty quantification (UQ) in imaging is that this error bar is actionable — that

a clinician, or an automated triage rule, can use it to decide whether to trust a value, escalate to another modality, or spare a patient an intervention. That premise hides three distinct questions, and they must be answered in order:

1. **Trust.** Can the reported interval be believed — does it mean what it says?
2. **Value of information.** *Given* a trusted bar, is it worth anything to the decision the clinician is actually making?
3. **Action.** When the bar is deployed and acted upon, where does it break?

A research program is most honest when its separate results turn out to answer one question each, and chain. This perspective makes that chain explicit for a multi-project IVIM UQ program: a calibration ruler (the *trust* anchor), a decision-theoretic pricing of a trusted bar (the *value* anchor), and two independent accounts of where deployment fails (the *action* anchors). We add no new experiment. We argue a synthesis, and we let the synthesis fall where the evidence pushes it.

Where it pushes is onto a single parameter. IVIM models the diffusion-weighted signal as a biexponential, $S(b)/S_0 = f e^{-bD^*} + (1-f) e^{-bD}$, separating tissue diffusion D from a fast pseudo-diffusion compartment D^* attributed to microvascular perfusion. D^* is the parameter the program keeps hitting a wall on, in every one of the three questions. We use it as the spine of the argument and call the resulting thread “ D^* cross-modally orphaned”: the program’s whole apparatus works for D (and the perfusion fraction f) and collapses for D^* . Crucially, this is an *identifiability* statement, not a claim that UQ is futile. The distinction is the point of the paper.

2 Trust: the calibration ruler

Before an error bar can have decision value it must be trustworthy: the reported interval must have the coverage it advertises. The trust anchor (Fashion [1])^P builds a model-agnostic *calibration ruler* — a check of whether a reported IVIM uncertainty is honest — and shows that a skew-aware posterior (an MCMC quantile interval) restores *marginal* coverage where symmetric $\pm\sigma$ intervals are miscalibrated. In its retooled, boundary-railing-first form the ruler reports this as a scoped secondary finding under an honest Cramér–Rao baseline, with the load-bearing residual being a *conditional* failure concentrated in the high- D^* regime.

The reading that matters for the arc: trust is achievable for D and f , and for D^* only marginally — with a residual high- D^* hole. That hole is the seam the rest of the arc pulls on. (This claim is ^P: Fashion’s calibration behaviour must survive review; it is a *contextual* anchor, not part of the load-bearing spine of Sec. 5.)

3 Value of information: pricing a trusted bar

Suppose the bar can be trusted. Is it worth anything? A calibrated interval is not automatically decision-relevant; its value is whatever it changes about the action taken. The value anchor (Minos [2]) prices exactly this for a treat / spare / escalate decision, and — importantly for a paper built before publication — its *theory half* (the Plumline results) is machine-verified and publication-independent. Three results carry the section.

The decision–calibration gap. The interval scale that is optimal for the *decision* departs from the scale that is optimal for *coverage* whenever the posterior is skewed and the cost is asymmetric. To leading order the gap is $G = \frac{1}{6} |z^*(\lambda)| \gamma$, with skewness γ and an asymmetry term $z^*(\lambda)$ set by the cost ratio λ (Plumblin Theorem 1). A symmetric problem ($\gamma = 0$) or a symmetric cost ($\lambda = 1$) closes the gap; IVIM is neither.

The value of information is second-order. The gain from re-calibrating the bar to the decision, rather than to coverage, is $V = \text{EU}(\tau^*) - \text{EU}(\tau_{\text{stat}}) = \frac{1}{2} |\text{EU}''(\tau^*)| G^2 = O(\gamma^2)$ (Plumblin Proposition 3 — the *Delphi* result). The value is *quadratic* in the gap: a calibrated bar changes the decision’s value only above a curvature-set “worth-it” floor. This is a leading-order, small-gap statement and we scope it as such — it characterizes the regime near the decision boundary where the quadratic approximation holds, not arbitrarily large mis-scalings.

The label-free detectability floor. A label-free validity monitor can catch *observable* distribution drift but is provably blind ($\text{AUC} = \frac{1}{2}$) to a *hidden* channel that induces the same regret (Plumblin Theorem 2). This is the formal case that label-free monitoring cannot, alone, certify a deployed bar; labeled spot-checks are required.

The reading for the arc: a trusted, calibrated bar has real, quantifiable decision value — but only where the gap is appreciable, and a label-free monitor cannot certify it by itself. The *theory* here is SOLID (independent of publication); Minos’s *applied* half, which consumes the trust and action anchors, is ^P. No applied number from Minos is load-bearing for the spine below.

4 Action: where acting on the bar breaks

When the bar is deployed and acted on, two *independent* walls appear — one about the interval’s size, one about which parameter it is even possible to size.

Wrong size (Lethe). On real same-day scan–rescan data (ACRIN-6698 [3], $n \approx 76$), a conformal interval calibrated for accuracy is far too narrow to cover real test–retest variation of D (Lethe [4]) ^P. This is not a contradiction: an interval scaled for *measurement accuracy* is *expected* to under-cover *repeatability*, because accuracy coverage and repeatability coverage are different targets and cannot be conflated. Precision is not coverage. The honest-limitation verdict is sharply scoped and survives as a caution about *which* coverage a deployed interval actually provides.

Wrong parameter (Gauge): an identifiability wall. D^* is essentially unidentifiable per voxel (Gauge [5]). The argument is a Cramér–Rao one and we are careful to frame it as such: the Cramér–Rao lower bound (CRLB) is the smallest variance any unbiased estimator can achieve from the IVIM b -value curve at a given operating point — a statement about *local identifiability*, not an impossibility proof. Scoped to the segmented biexponential IVIM scheme and the SNR band 10–100, the CRLB on D^* grows by a factor of ~ 6 across the physiological D^* range (reproduced in this work, Fig. 1; reproduced factor $5.7\times$), and in the high- D^* tercile the CRLB standard deviation reaches $\sim 1.12\times$ the width of that tercile — the regime is, in this acquisition, unresolvable. That conformal interval widths track the CRLB ($r = 0.77$) confirms the intervals are honest about this: they widen exactly where identifiability fails. What cannot be guaranteed label-free is coverage *conditional* on a latent axis the data do not identify — the IVIM instance of the impossibility of distribution-free conditional coverage [6], and a concrete instantiation of the hidden, undetectable channel of Sec. 3.




Figure 1: The pseudo-diffusion identifiability wall, reproduced. Cramér–Rao lower bound on D^* (standard deviation, S_0 jointly estimated) versus true D^* across the SNR band, for the segmented biexponential IVIM scheme. The bound grows $\sim 6\times$ across the D^* range and, in the shaded high- D^* tercile, reaches the scale of the tercile width itself: the regime is unresolvable per voxel. This is an identifiability (CRLB) limit scoped to this acquisition and SNR band, not an impossibility proof. Self-contained re-derivation (`scripts/crlb_wall.py`); cross-checked against the committed Gauge anchor.

The reading for the arc: acting on the bar is safe for the parameters the program can identify and size; it is *not* safe for D^* , which sits behind an identifiability wall a label-free monitor cannot see — so labeled repeatability spot-checks are mandatory, not optional.

5 The thread: D^* is cross-modally orphaned

D^* is where trust, value, and action fail *together*, and the failure extends beyond IVIM, across modalities. Three independent observations form the spine. The first two are **in-repository, reproduced anchors** (the load-bearing core); the third is corroborating external literature.

(1) Unidentifiable from its own signal. D^* cannot be recovered per voxel from the IVIM b -value curve: the CRLB grows $\sim 6\times$ in the high- D^* regime and reaches the regime-bin scale there (Sec. 4, Fig. 1). This is reproduced from first principles in this work, not merely cited.

(2) Un-scalable to its own repeatability. On the real same-day test-retest arm, the D^* interval width does *not* track scan-rescan scatter: Spearman $r = -0.17$ ($p = 0.13$, $n = 76$; 95% BCa bootstrap CI $[-0.39, +0.05]$, *spanning zero*) — a null. By contrast the well-identified D tracks its own scatter at $r = +0.60$ (95% CI $[+0.42, +0.72]$, strictly positive). The D^* null is reported *with its interval*, not as a bare point estimate, precisely because a null is only meaningful with its uncertainty; the CI’s spanning zero is the load-bearing fact. We stress the caveat that this is a region-level, whole-tumor-ROI repeatability-tracking association on unregistered same-day repeats

— not an in-vivo coverage, accuracy, or calibration claim. The D^* negative is *evidence for* the identifiability thesis, not a gap in it.

(3) Weakly and inconsistently anchored to an independent perfusion measurement.

If D^* measured perfusion, it should correlate with an independent perfusion readout. The cross-modal correlation of D^* with DCE-MRI K^{trans} is, at best, weak and cohort-dependent: $r = 0.389$ ($p < 0.001$) in one rectal-cancer cohort [7] but *non-significant* in another [8], which also reports only moderate D^* reproducibility (ICC = 0.55). We frame this honestly as *weak and cohort-inconsistent* — significant in one cohort, null in the other — *not* uniformly “non-significant”; both are cited. Tellingly, the composite $f \cdot D^*$ correlates better ($r = 0.533$), consistent with f — not D^* — carrying the recoverable perfusion signal. This strand is corroborating; the load-bearing spine is (1)–(2).

The synthesis. D^* is orphaned: a trusted ruler, a priced decision, and a sized interval all work for D and f and collapse for D^* — unidentifiable from its own signal, un-scalable to its own repeatability, and only weakly and inconsistently tied to an independent perfusion readout.

6 End-stage caution

The defensible end-stage statement is not that uncertainty quantification fails. It is the sharper claim the program’s own machinery earns: *the value of a calibrated error bar is real and quantifiable exactly where the parameter is identifiable*, and D^* marks the identifiability wall where trust, value, and action all terminate. Operationally this is a triage rule, and we are careful not to overclaim it at the boundary: trust the well-identified diffusion map D ; read high- D^* as unresolved rather than as a measured perfusion value; price the bar’s decision value only where the decision–calibration gap is appreciable (Sec. 3); and, because the failing axis is one no label-free monitor can see, treat labeled repeatability spot-checks as mandatory. None of this is a claim that D^* “works” once these cautions are applied — the spine is an honest negative, and we let it stand as one.

7 Provisional discipline and submission hold

This synthesis was assembled *before* its in-program anchors published. We treat that as a discipline, not a disclaimer. Every claim that depends on an unpublished anchor is marked ^P. Every external number traces to a verified primary source. The load-bearing spine of Sec. 5 — the D^* identifiability wall and the D^* test–retest null with its bootstrap CI — rests on in-repository, reproduced anchors and is *self-contained*: the trust (Fashion) and value (Minos) projects are *contextual* forward-citations, and the core argument survives even if their final numbers shift in review. We have confirmed this explicitly: no held or provisional number from those anchors enters the spine.

Accordingly, the manuscript is complete and reproduces green against its in-repository anchors, but **submission is held** behind an explicit release gate until the trust and value anchors publish (and ideally the action anchors as well). Reproduction and release are separate concerns: `reproduce.sh` regenerates every in-repository anchor; the release gate (`release_gate.py`, `submit.sh`) refuses to submit while either anchor is unpublished. The path to lift the hold, and the per-number re-check that fires when each anchor is accepted, are documented in `SUBMISSION_BLOCK.md` and `PROVISIONAL_LEDGER.md`.

Data and code availability

This is a synthesis; it generates no new data. The in-repository anchors are reproduced by `bash reproduce.sh`: a self-contained re-derivation of the CRLB identifiability wall (Fig. 1; `scripts/crlb_wall.py`), the D^* test-retest correlation interval (`scripts/retest_ci.py`, carrying the seeded BCa bootstrap CI and reproducing its Fisher- z twin), and the external D^*-K^{trans} evidence table (`scripts/dstar_ktrans.py`). The underlying test-retest data are the open ACRIN-6698 / I-SPY2 breast DWI cohort (TCIA, CC-BY-4.0 [3]). No private clinical data are touched. All manuscript numbers are auto-generated into `numbers.tex` by `paper/consistency.py` and trace to committed sources.

Forward-citation note. Fashion [1] and Minos [2] are cited as in-program forward-references and are unpublished at the time of writing; Gauge [5] and Lethe [4] are companion manuscripts. No published DOI is asserted for an unpublished anchor; each is swapped to its published reference at finalization per the checklist in `PROVISIONAL_LEDGER.md`.

References

- [1] Avery Karlin. Fashion: a boundary-railing calibration ruler for IVIM uncertainty (the trust anchor). In review, NMR in Biomedicine (retooled). Code archive: Zenodo 10.5281/zenodo.20649669, 2026.
- [2] Avery Karlin. Minos: the value of information of decision-calibration in IVIM (plumbline theory: gap law, second-order VoI, label-free detectability floor). Provisional; theory half machine-verified and publication-independent, 2026.
- [3] David C. Newitt, Savannah C. Partridge, Zheng Zhang, et al. ACRIN 6698/I-SPY2 breast DWI (diffusion-weighted MRI). The Cancer Imaging Archive (TCIA), 2021. CC-BY-4.0.
- [4] Avery Karlin. Lethe: precision is not coverage – the wrong-size limit of conformal IVIM intervals on real test-retest data. Constrained-validation verdict rendered; re-pointed to NMR in Biomedicine, 2026.
- [5] Avery Karlin. Distribution-free conformal coverage for IVIM parameter maps, and the identifiability wall in the pseudo-diffusion compartment. Manuscript assembled; target Magnetic Resonance in Medicine, 2026.
- [6] Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. The limits of distribution-free conditional predictive inference. *Information and Inference*, 10(2):455–482, 2021. doi: 10.1093/imaiai/iaaa017.
- [7] Hongliang Sun, Yanqing Xu, Qing Xu, Jianghui Duan, Hua Zhang, Ting Liu, Lufan Li, Queenie Chan, Sheng Xie, and Wu Wang. Correlation between intravoxel incoherent motion and dynamic contrast-enhanced magnetic resonance imaging parameters in rectal cancer. *Academic Radiology*, 26(7):e134–e141, 2019. doi: 10.1016/j.acra.2018.08.012. PMID: 30268719.
- [8] Xiaodan Yang, Xiao Xiao, Bei Lu, Yan Chen, Zhibo Wen, and Shenping Yu. Perfusion-sensitive parameters of intravoxel incoherent motion MRI in rectal cancer: evaluation of reproducibility and correlation with dynamic contrast-enhanced MRI. *Acta Radiologica*, 60(5):569–577, 2019. doi: 10.1177/0284185118791201. PMID: 30114928.